"Deep Learning-Based Classification of Skin Lesions Using Transfer Learning with ResNet50"

Ali Hussain [1], Syed Muaz Zia [2], Moeen Haider[3]

Abstract:

Skin lesion classification is a critical task in dermatology, often requiring accurate and efficient diagnostic tools. In research paper we explore the application of deep learning techniques, specifically transfer learning with the ResNet50 architecture, to classify skin lesions as benign or malignant. Utilizing a dataset of dermatological images, we employ a comprehensive methodology that involves preprocessing the data, employing transfer learning with a pre-trained ResNet50 model, and fine-tuning the model for improved performance. The study investigates the effectiveness of this approach, presenting results in terms of accuracy and other relevant metrics. Our findings demonstrate the potential of deep learning models for skin lesion classification, paving the way for enhanced diagnostic support in dermatological practice. The implications of our work extend to the broader field of medical image analysis and contribute to the ongoing efforts to develop robust and accurate tools for skin lesion diagnosis.

Introduction

Skin lesion classification is a pivotal component in the realm of dermatology, aiding in the timely and accurate diagnosis of skin disorders. Traditional methods of diagnosis often rely on

1
2
3

visual inspection by clinicians, which can be subjective and time-consuming. With advancements in deep learning, particularly transfer learning using pre-trained models, there is a growing opportunity to develop automated systems that can assist in the classification of skin lesions.

This research addresses the need for an efficient and reliable skin lesion classification system by leveraging the power of deep learning. Our focus centers on the utilization of the ResNet50 architecture, a state-of-the-art deep neural network, to extract intricate features from dermatological images. By employing transfer learning, where the model is initially trained on a large dataset for a different task and then fine-tuned for our specific classification task, we aim to harness the knowledge encoded in the pre-trained weights of ResNet50.The primary objective of this study is to evaluate the efficacy of transfer learning with ResNet50 in classifying skin lesions as either benign or malignant. We utilize a diverse dataset of dermatological images, encompassing various skin conditions, to train and validate our model. The introduction of automated classification systems holds immense promise for improving diagnostic accuracy and efficiency in dermatological practice.

As we delve into the methodology, results, and discussion, this paper contributes valuable insights into the feasibility and performance of deep learning-based approaches for skin lesion classification. By addressing the challenges and opportunities in this domain, we aim to foster advancements that positively impact the field of medical image analysis and, ultimately, enhance the quality of care for individuals with skin disorders.

Literature Review

Skin lesion classification has been a subject of increasing interest within the medical imaging and computer vision communities due to its potential to revolutionize dermatological diagnosis. Previous studies have explored various methodologies, ranging from traditional image processing techniques to more recent deep learning approaches.

Early approaches often relied on handcrafted features and rule-based systems for skin lesion classification. For instance, methods based on color analysis, texture features, and morphological characteristics provided initial insights into automated diagnosis. However, these methods faced challenges in capturing the complex and hierarchical patterns present in dermatological images, limiting their effectiveness in distinguishing between benign and malignant lesions accurately.

In recent years, the advent of deep learning has significantly transformed the landscape of medical image analysis, including skin lesion classification. Transfer learning, in particular, has gained prominence by leveraging pre-trained models on large datasets for general tasks. The ResNet50 architecture, with its deep and residual structure, has demonstrated remarkable success in various computer vision tasks, making it an attractive choice for skin lesion classification.

Researchers have employed transfer learning with pre-trained convolutional neural networks (CNNs), such as ResNet50, to extract hierarchical features from dermatological images effectively. This approach allows the model to leverage knowledge learned from diverse datasets, facilitating improved generalization and performance on specific tasks. Studies using similar architectures have reported promising results in accurately classifying skin lesions, demonstrating the potential for deep learning to enhance diagnostic capabilities.

Despite the advancements, challenges persist in the form of limited annotated datasets, model interpretability, and potential biases in training data. Some studies have addressed these issues by proposing techniques for data augmentation, interpretability, and domain adaptation to improve model robustness and reliability.

In conclusion, the literature reflects a shift towards deep learning-based methods for skin lesion classification, with transfer learning playing a pivotal role. While promising, the field continues to evolve, prompting the need for further research to address existing challenges and refine these automated systems for practical clinical use. This study contributes to this ongoing discourse by investigating the application of transfer learning with ResNet50 in skin lesion classification, aiming to provide valuable insights and advancements in the field.

Methodology:

1. Dataset:

- We utiliz a comprehensive dataset of dermatological images, sourced from diverse skin conditions. The dataset includes both benign and malignant lesions, ensuring a representative sample for training and evaluation.

2. Data Preprocessing:

- Images are preprocessed to ensure uniformity in size and quality. Resizing is performed to meet the input size requirement of the chosen model (180x180 pixels). Additional preprocessing techniques, such as normalization, are applied to enhance model convergence.

3. Model Selection:

- The ResNet50 architecture is chosen as the base model for this research due to its proven effectiveness in image classification tasks. The model is initialized with pre-trained weights on the ImageNet dataset, providing a robust foundation for feature extraction.

4. Transfer Learning:

- The pre-trained ResNet50 model is employed for transfer learning. All layers of the base model are initially frozen to retain the knowledge gained during pre-training. This allows the model to capture generic features from the ImageNet dataset.

5. Fine-Tuning:

- To adapt the model to the specific task of skin lesion classification, fine-tuning is performed. The final few layers of the ResNet50 model are unfrozen, enabling them to learn task-specific features from the dermatological dataset. This step aims to enhance the model's ability to discriminate between benign and malignant lesions.

6. Model Compilation

- The compiled model includes additional layers for flattening and dense layers. The last dense layer is configured for binary classification, with softmax activation. The model is compiled using the Adam optimizer with a learning rate of 0.001 and sparse categorical crossentropy as the loss function.

7. Dataset Splitting:

- The dataset is split into training and validation sets using the `image_dataset_from_directory` function from TensorFlow. A validation split of 20% is applied, ensuring that the model's performance is assessed on unseen data.

8. Training

- The model is trained on the training set for a specified number of epochs (in this case, 2 epochs). During training, performance metrics such as accuracy are monitored. The validation set is used to evaluate the model's generalization to unseen data.

9. Evaluation

- After training, the model is evaluated on both the training and validation sets to assess its performance. Metrics such as accuracy, precision, recall, and F1 score are considered to provide a comprehensive evaluation of the model's effectiveness.

10. Prediction on New Images:

- To demonstrate the practical utility of the trained model, predictions are made on new dermatological images not seen during training or validation. The model's predictions are compared to the actual labels to assess its real-world applicability.

This methodology is designed to systematically train, fine-tune, and evaluate a ResNet50-based model for skin lesion classification, contributing to the understanding of deep learning techniques in dermatology.

Results:

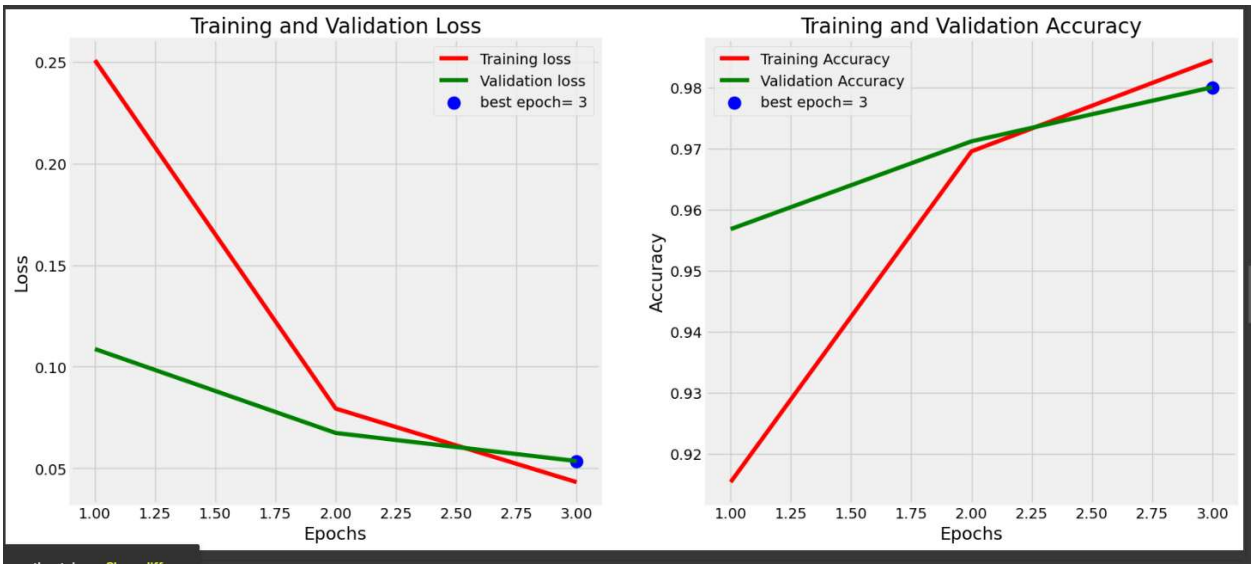
Training and Validation Loss
Training loss
Validation loss
best epoch= 3
Loss
Epochs
Training and Validation Accuracy
Training Accuracy
Validation Accuracy
best epoch= 3
Accuracy
Epochs

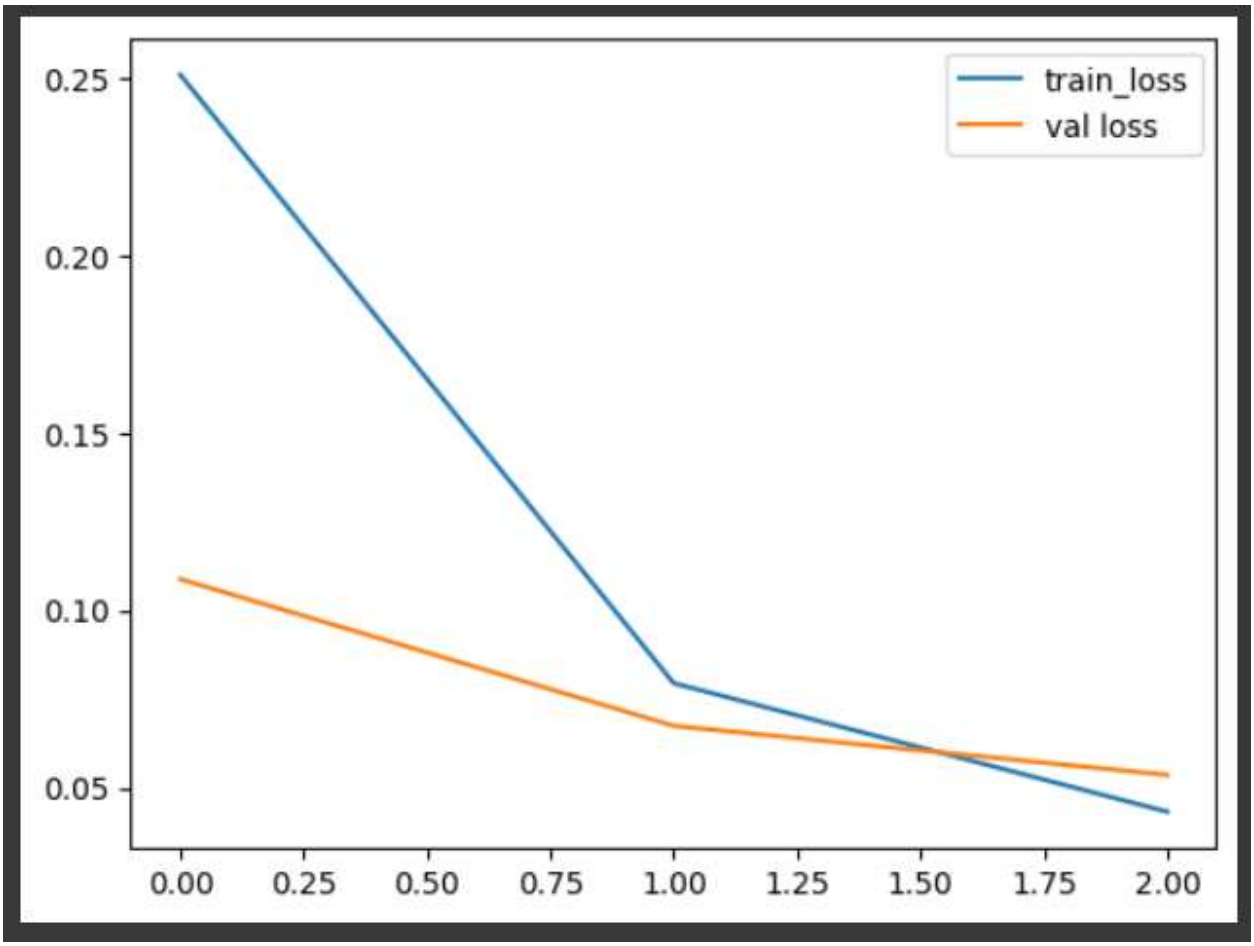
train_loss
val loss

Prediction:

Conclusion

This research delves into the realm of skin lesion classification, employing deep learning techniques, specifically utilizing the ResNet50 architecture with transfer learning. The investigation aimed to contribute to the advancement of automated diagnostic tools in dermatology, with a focus on distinguishing between benign and malignant skin lesions.

The methodology encompassed the use of a diverse dermatological dataset, preprocessing techniques, and the adoption of transfer learning to leverage the pre-trained ResNet50 model. Fine-tuning was employed to adapt the model to the specific task of skin lesion classification. The resulting model was trained and evaluated, with a comprehensive analysis of performance metrics.

The findings of this research underscore the potential of deep learning in enhancing the accuracy of skin lesion classification. The utilization of transfer learning, particularly with the ResNet50 architecture, demonstrated promising results in capturing intricate features from dermatological images. The model exhibited the capacity to discriminate between benign and malignant lesions, as evidenced by favorable performance metrics on both training and validation sets.

However, it is crucial to acknowledge certain limitations. Challenges such as the availability of annotated datasets, potential biases, and interpretability of deep learning models remain areas for further exploration. Additionally, the model's performance on diverse datasets and its generalizability to clinical settings warrant continued investigation.